

```
In [ ]: import numpy as np
import pandas as pd
from numpy import unique, argmax
from tensorflow.keras.datasets.mnist import load_data
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D
from tensorflow.keras.layers import Dense
from tensorflow.keras.layers import Flatten
from tensorflow.keras.layers import Dropout
from tensorflow.keras.utils import plot_model
import matplotlib.pyplot as plt
from tensorflow.keras.datasets import mnist

In [ ]: (train_x, train_y), (test_x, test_y) = mnist.load_data()

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 [=====] - 1s 0us/step

In [ ]: print(train_x.shape, train_y.shape)
print(test_x.shape, test_y.shape)

(60000, 28, 28) (60000,)
(10000, 28, 28) (10000,)

In [ ]: train_x = train_x.astype('float32')/255.0
test_x = test_x.astype('float32')/255.0

In [ ]: fig = plt.figure(figsize = (20,5))
for i in range(20):
    ax=fig.add_subplot(2, 10, i+1, xticks=[], yticks=[])
    ax.imshow(np.squeeze(train_x[i]), cmap='gray')
    ax.set_title(train_y[i])

5      0      4      1      9      2      1      3      1      4
3      5      3      6      1      7      2      8      6      9

In [ ]: shape = train_x.shape[1:]
shape

Out[ ]: (28, 28)

In [ ]: from tensorflow.keras.layers import MaxPooling2D as MaxPool2D
model = Sequential()
#adding convolutional layer
model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
model.add(MaxPool2D((2,2)))
model.add(Conv2D(48, (3,3), activation='relu'))
model.add(MaxPool2D((2,2)))

In [ ]: model.add(Dropout(0.5))
model.add(Flatten())
model.add(Dense(500, activation='relu'))
model.add(Dense(10, activation='softmax'))

In [ ]: model.summary()

Model: "sequential"

Layer (type)                Output Shape                Param #
=====
conv2d (Conv2D)              (None, 26, 26, 32)         320

max_pooling2d (MaxPooling2D) (None, 13, 13, 32)         0

conv2d_1 (Conv2D)            (None, 11, 11, 48)         13872

max_pooling2d_1 (MaxPooling2D) (None, 5, 5, 48)         0

dropout (Dropout)            (None, 5, 5, 48)           0

flatten (Flatten)             (None, 1200)                0

dense (Dense)                 (None, 500)                 600500

dense_1 (Dense)               (None, 10)                  5010

=====
Total params: 619,702
Trainable params: 619,702
Non-trainable params: 0

In [ ]: model.compile(optimizer='adam', loss = 'sparse_categorical_crossentropy',metrics= ['accuracy'])
x=model.fit(train_x, train_y, epochs=10, batch_size = 128, verbose= 2 , validation_split = 0.1)

Epoch 1/10
422/422 - 14s - loss: 0.2410 - accuracy: 0.9250 - val_loss: 0.0598 - val_accuracy: 0.9833 - 14s/epoch - 33ms/step
Epoch 2/10
422/422 - 2s - loss: 0.0809 - accuracy: 0.9746 - val_loss: 0.0405 - val_accuracy: 0.9890 - 2s/epoch - 4ms/step
Epoch 3/10
422/422 - 2s - loss: 0.0603 - accuracy: 0.9806 - val_loss: 0.0371 - val_accuracy: 0.9898 - 2s/epoch - 4ms/step
Epoch 4/10
422/422 - 2s - loss: 0.0481 - accuracy: 0.9847 - val_loss: 0.0329 - val_accuracy: 0.9913 - 2s/epoch - 4ms/step
Epoch 5/10
422/422 - 2s - loss: 0.0405 - accuracy: 0.9866 - val_loss: 0.0303 - val_accuracy: 0.9913 - 2s/epoch - 4ms/step
Epoch 6/10
422/422 - 2s - loss: 0.0360 - accuracy: 0.9880 - val_loss: 0.0269 - val_accuracy: 0.9928 - 2s/epoch - 5ms/step
Epoch 7/10
422/422 - 2s - loss: 0.0331 - accuracy: 0.9893 - val_loss: 0.0302 - val_accuracy: 0.9915 - 2s/epoch - 4ms/step
Epoch 8/10
422/422 - 2s - loss: 0.0286 - accuracy: 0.9906 - val_loss: 0.0278 - val_accuracy: 0.9933 - 2s/epoch - 4ms/step
Epoch 9/10
422/422 - 2s - loss: 0.0276 - accuracy: 0.9907 - val_loss: 0.0306 - val_accuracy: 0.9925 - 2s/epoch - 4ms/step
Epoch 10/10
422/422 - 2s - loss: 0.0232 - accuracy: 0.9926 - val_loss: 0.0277 - val_accuracy: 0.9937 - 2s/epoch - 4ms/step

In [ ]: loss, accuracy= model.evaluate(test_x, test_y, verbose = 0)
print(f'Accuracy: {accuracy*100}%')

Accuracy: 99.32000041007996

In [14]: model.save(r'C:\Users\92341\Desktop\DipLab\final_model.h5')

In [29]: from numpy import argmax
from tensorflow.keras.utils import load_img
from tensorflow.keras.utils import img_to_array
from keras.models import load_model
import matplotlib.pyplot as plt
def load_image(filename):
    img = load_img(filename, grayscale=True, target_size=(28, 28))
    plt.figure(figsize=(5,5))
    plt.imshow(img, cmap='gray')
    plt.show()
    img = img_to_array(img)
    img = img.reshape(1, 28, 28, 1)
    img = img.astype('float32')
    img = img / 255.0
    return img

In [36]: def run_example():
img = load_image(r'./content/Screenshot_2023-07-20_055517.png')
model = load_model(r'C:\Users\92341\Desktop\DipLab\final_model.h5')
predict_value = model.predict(img)
digit = argmax(predict_value)
return 'Predicted',digit
run_example()

0
5
10
15
20
25
0
5
10
15
20
25

1/1 [=====] - 0s 160ms/step
('Predicted', 7)

Out[36]:

In [37]: import numpy as np
import pandas as pd
from numpy import unique, argmax
from tensorflow.keras.datasets.mnist import load_data
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D
from tensorflow.keras.layers import Dense
from tensorflow.keras.layers import Flatten
from tensorflow.keras.layers import Dropout
from tensorflow.keras.utils import plot_model
import matplotlib.pyplot as plt
from tensorflow.keras.datasets import mnist

In [38]: (train_x, train_y), (test_x, test_y) = mnist.load_data()

In [39]: print(train_x.shape, train_y.shape)
print(test_x.shape, test_y.shape)

(60000, 28, 28) (60000,)
(10000, 28, 28) (10000,)

In [40]: train_x = train_x.astype('float32')/255.0
test_x = test_x.astype('float32')/255.0

In [41]: fig = plt.figure(figsize = (20,5))
for i in range(20):
    ax=fig.add_subplot(2, 10, i+1, xticks=[], yticks=[])
    ax.imshow(np.squeeze(train_x[i]), cmap='gray')
    ax.set_title(train_y[i])

5      0      4      1      9      2      1      3      1      4
3      5      3      6      1      7      2      8      6      9

In [42]: shape = train_x.shape[1:]
shape

Out[42]: (28, 28)

In [43]: from tensorflow.keras.layers import MaxPooling2D as MaxPool2D
model = Sequential()
#adding convolutional layer
model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
model.add(MaxPool2D((2,2)))
model.add(Conv2D(48, (3,3), activation='relu'))
model.add(MaxPool2D((2,2)))

In [44]: model.add(Dropout(0.5))
model.add(Flatten())
model.add(Dense(500, activation='relu'))
model.add(Dense(10, activation='softmax'))

In [45]: model.summary()

Model: "sequential_1"

Layer (type)                Output Shape                Param #
=====
conv2d_2 (Conv2D)              (None, 26, 26, 32)         320

max_pooling2d_2 (MaxPooling2D) (None, 13, 13, 32)         0

conv2d_3 (Conv2D)            (None, 11, 11, 48)         13872

max_pooling2d_3 (MaxPooling2D) (None, 5, 5, 48)         0

dropout_1 (Dropout)            (None, 5, 5, 48)           0

flatten_1 (Flatten)             (None, 1200)                0

dense_2 (Dense)                 (None, 500)                 600500

dense_3 (Dense)               (None, 10)                  5010

=====
Total params: 619,702
Trainable params: 619,702
Non-trainable params: 0

In [46]: model.compile(optimizer='adam', loss = 'sparse_categorical_crossentropy',metrics= ['accuracy'])
x=model.fit(train_x, train_y, epochs=10, batch_size = 128, verbose= 2 , validation_split = 0.1)

Epoch 1/10
422/422 - 6s - loss: 0.2309 - accuracy: 0.9275 - val_loss: 0.0619 - val_accuracy: 0.9825 - 6s/epoch - 14ms/step
Epoch 2/10
422/422 - 2s - loss: 0.0783 - accuracy: 0.9749 - val_loss: 0.0456 - val_accuracy: 0.9868 - 2s/epoch - 4ms/step
Epoch 3/10
422/422 - 2s - loss: 0.0590 - accuracy: 0.9822 - val_loss: 0.0339 - val_accuracy: 0.9910 - 2s/epoch - 4ms/step
Epoch 4/10
422/422 - 2s - loss: 0.0489 - accuracy: 0.9841 - val_loss: 0.0377 - val_accuracy: 0.9898 - 2s/epoch - 4ms/step
Epoch 5/10
422/422 - 2s - loss: 0.0427 - accuracy: 0.9861 - val_loss: 0.0302 - val_accuracy: 0.9915 - 2s/epoch - 4ms/step
Epoch 6/10
422/422 - 2s - loss: 0.0343 - accuracy: 0.9891 - val_loss: 0.0260 - val_accuracy: 0.9932 - 2s/epoch - 5ms/step
Epoch 7/10
422/422 - 2s - loss: 0.0304 - accuracy: 0.9903 - val_loss: 0.0275 - val_accuracy: 0.9923 - 2s/epoch - 5ms/step
Epoch 8/10
422/422 - 2s - loss: 0.0280 - accuracy: 0.9910 - val_loss: 0.0267 - val_accuracy: 0.9930 - 2s/epoch - 4ms/step
Epoch 9/10
422/422 - 2s - loss: 0.0250 - accuracy: 0.9921 - val_loss: 0.0254 - val_accuracy: 0.9935 - 2s/epoch - 4ms/step
Epoch 10/10
422/422 - 2s - loss: 0.0244 - accuracy: 0.9916 - val_loss: 0.0245 - val_accuracy: 0.9938 - 2s/epoch - 4ms/step

In [47]: loss, accuracy= model.evaluate(test_x, test_y, verbose = 0)
print(f'Accuracy: {accuracy*100}%')

Accuracy: 99.2799973297119

In [48]: model.save(r'C:\Users\92341\Desktop\DipLab\final_model.h5')

In [57]: from numpy import argmax
from tensorflow.keras.utils import load_img
from tensorflow.keras.utils import img_to_array
from keras.models import load_model
import matplotlib.pyplot as plt
# load and prepare the image
def load_image(filename):
    img = load_img(filename, grayscale=True, target_size=(28, 28))
    plt.figure(figsize=(5,5))
    plt.imshow(img, cmap='gray')
    plt.show()
    # convert to array
    img = img_to_array(img)
    # reshape into a single sample with 1 channel
    img = img.reshape(1, 28, 28, 1)
    # prepare pixel data
    img = img.astype('float32')
    img = img / 255.0
    return img
# load an image and predict the class
def run_example():
    # load the image
    img = load_image(r'./content/Screenshot_2023-07-20_055517.png')
    # load model
    model = load_model(r'C:\Users\92341\Desktop\DipLab\final_model.h5')
    # predict the class
    predict_value = model.predict(img)
    digit = argmax(predict_value)
    print('Predicted',digit)
    # entry point, run the example
    run_example()

/usr/local/lib/python3.10/dist-packages/keras/utils/image_utils.py:409: UserWarning: grayscale is deprecated. Please use color_mode = "grayscale"
warnings.warn(

0
5
10
15
20
25
0
5
10
15
20
25

1/1 [=====] - 0s 74ms/step
Predicted 7

In [ ]:
```